


Part 1

Write your own KNN code without using any pack

```
myKnn = function(traindata, testdata, Ytrain, k = 1){  
  Ytest = array(0,dim=nrow(testdata))  
  
  for (row in 1:nrow(testdata)) {  
    distances = sqrt(rowSums((traindata - testdata[row,])^2))
```



As shown on the next page, the student's KNN result is not correct. This is because the calculation for “distances” is not correct.

```
traindata = matrix(1:8, 4, 2)  
testdata = matrix(1:4, 2, 2)  
traindata  
testdata[1, 3]  
traindata - testdata[1,]  
  
# instead one should use  
t(t(traindata) - testdata[1,])
```

Test your code with mydata when $K = 1, 3, 5$; compare your results with from the R command knn.

```
library(class)
test.pred = knn(traindata, testdata, Ytrain, k = 1)
table(Ytest, test.pred)
```

```
##      test.pred
## Ytest    0    1
##      0 3589 1411
##      1 1334 3666
```

```
test.pred = myKnn(traindata, testdata, Ytrain, k = 1)
table(Ytest, test.pred)
```

```
##      test.pred
## Ytest    0    1
##      0 1179 3821
##      1  973 4027
```

```
test.pred = knn(traindata, testdata, Ytrain, k = 3)
table(Ytest, test.pred)
```

```
##      test.pred
## Ytest    0    1
##      0 3642 1358
##      1 1076 3924
```

This assignment **DOES NOT** ask students to just **DUMP** their results in the report and leave us to **GUESS** what their results mean.

Apparently the results do not match, and the student should say something in the report.

```

library(class)
calc_error_rate <- function(predicted.value, true.value){
  return(mean(true.value!=predicted.value))
}

cvKNNaveErrorRate = function(K, traindata, Ytrain) {
  foldNum = 10
  n = nrow(traindata)
  foldSize = floor(n/foldNum)
  K = 3
  error = 0
  myIndex = sample(1 : n)

  for(runId in 1:foldNum){
    testSetIndex = ((runId-1)*foldSize + 1):(ifelse(runId == foldNum, n, runId
    testSetIndex = myIndex[testSetIndex]
    trainX = traindata[-testSetIndex, ]

```

No wonder the box plot (on the next page) looks so weird.


```

cvKNN = function(data) {
  foldNum = 10
  n = nrow(data$traindata)
  foldSize = floor(n/foldNum)
  KVector = seq(1, (nrow(data$traindata) - foldSize), 1)
  cvErrorRates = sapply(KVector, cvKNNaveErrorRate, data$traindata,
  result = list()

  bestK = max(KVector[cvErrorRates == min(cvErrorRates)])
  result$bestK = bestK
  result$train_error = cvErrorRates[KVector == bestK]

```

This is the CV error NOT the training error. Run knn with bestK on the training data, and compute the corresponding error.



Rotate X-tick label 45 or 60 degree

```

for (i in seq(1, length(k))){
  Ytest_pred = get_assignment(k[i], get_distance(traindata, testdata))
  print(paste("For K =", k[i]))
  print(table(Ytest, Ytest_pred))
  test.pred = knn(traindata, testdata, Ytrain, k = k[i], prob=TRUE)
  print(table(Ytest, test.pred))
  mm = which(attributes(test.pred)$prob == 0.5)
  if (length(mm)>0){
    print(paste(length(mm), "mismatches due to R's kNN distance buffer:"))
    print(mm)
  }
}

```

```

## [1] "For K = 5"
##      Ytest_pred
## Ytest    0    1
##      0 4067  933
##      1  687 4313
##      test.pred
## Ytest    0    1
##      0 4066  934
##      1  687 4313
## [1] "2 mismatches due to R's kNN distance buffer:"
## [1] 1419 5491

```

Students need to provide

- 1) IDs whose test.pred != Ytest_pred**
- 2) For IDs in 1), whose prob from knn are 0.5**

Are there two mismatches in the table shown above? Or you mean 2 potential mismatches...

Class implementation K=5:

Student 3

Hide

```
table(Ytest, test.pred.k5)
```

```
test.pred.k5
Ytest      0      1
0 2939 2061
1 1321 3679
```

Mismatch due to 0.5 probabilities

Why? How do you know it's due to 0.5 prob?

Hide

```
attributes(test.pred.k5)$prob[attributes(test.pred.k5)$prob == 0.5]
```

```
[1] 0.5 0.5
```

Even if a prediction prob = 0.5, the output from knn could still match with yours.

You have to show that the prediction prob (from knn) of that mismatched test point is equal to 0.5

```

csize = 10
p = 2
s = np.sqrt(float(1)/5)
n = 5000 # training sample for each class
N = 100

```

Wrong values for n and N

```

cv_err = np.array(cv_err)
best_k_idx = np.argmin(cv_err.mean(axis=1))
best_k = k_range[best_k_idx]

print('selected k = ', best_k)
print(f'the mean error of selected k: {cv_err[best_k_idx].mean()}')
print(f'the std error of selected k: {cv_err[best_k_idx].std()}')

```

```

selected k = 25
the mean error of selected k: 0.07
the std error of selected k: 0.045825756949558406

```

```


train_error_knn = []
test_error_knn = []

for r in range(50):
    X_train, Y_train, X_test, Y_test = generate_sim_data(center_0, cent
    Y_train_pred = knn_pred(X_train, X_train, Y_train, best_k)
    Y_test_pred = knn_pred(X_train, X_test, Y_train, best_k)

```

Each procedure should be evaluated on the **same** 50 sets of training/test data


Wrong CVkNN procedure:
You need to select best_k for each r

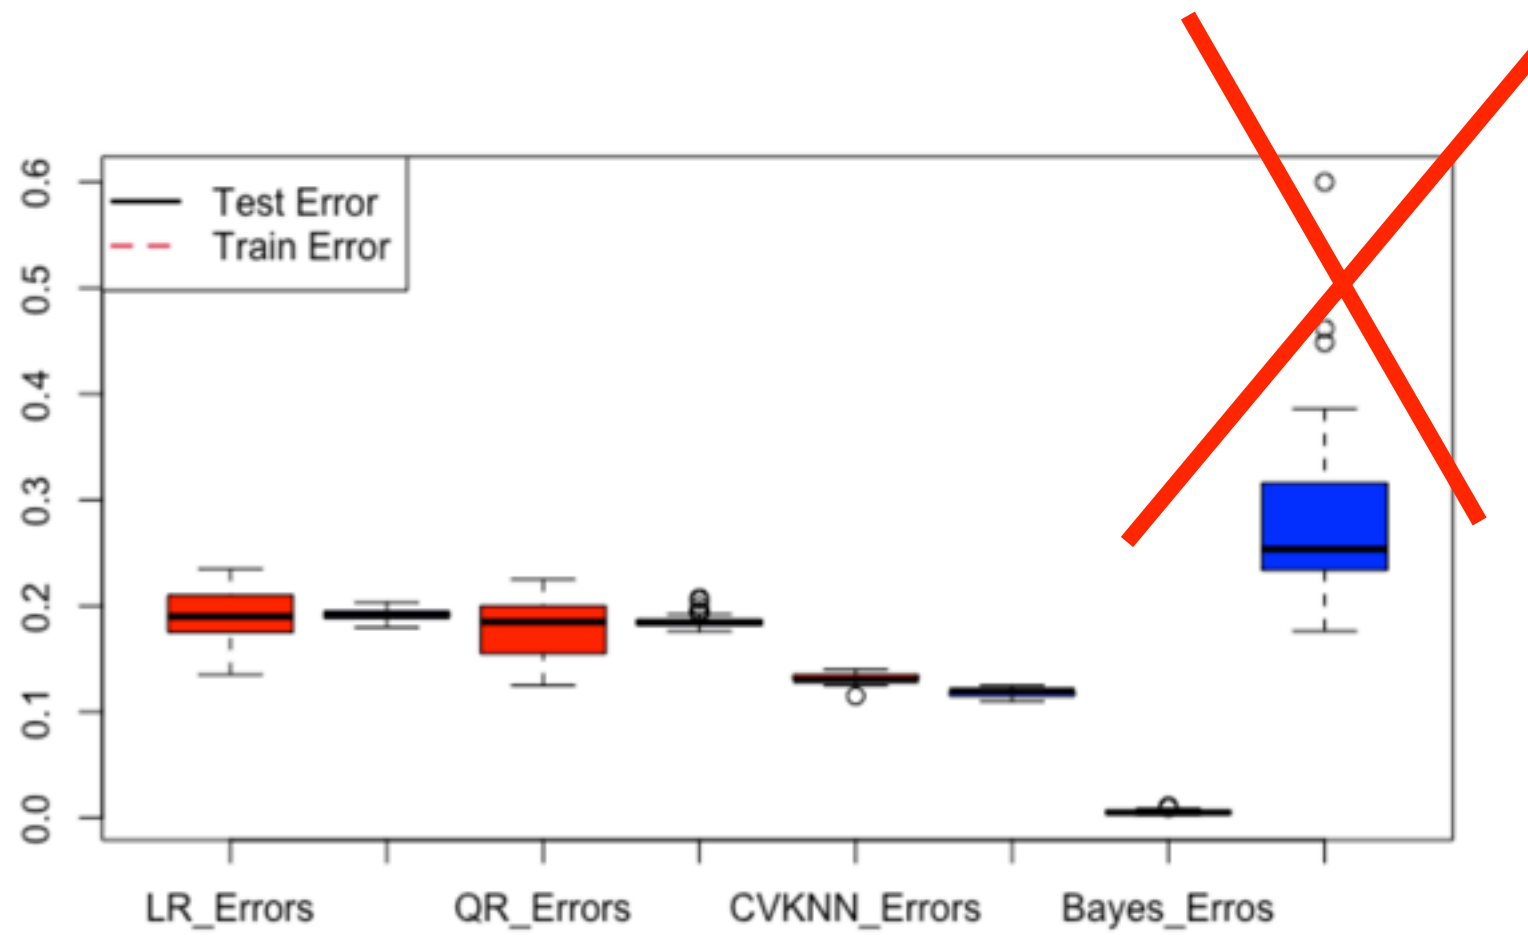


Apparently something is wrong.

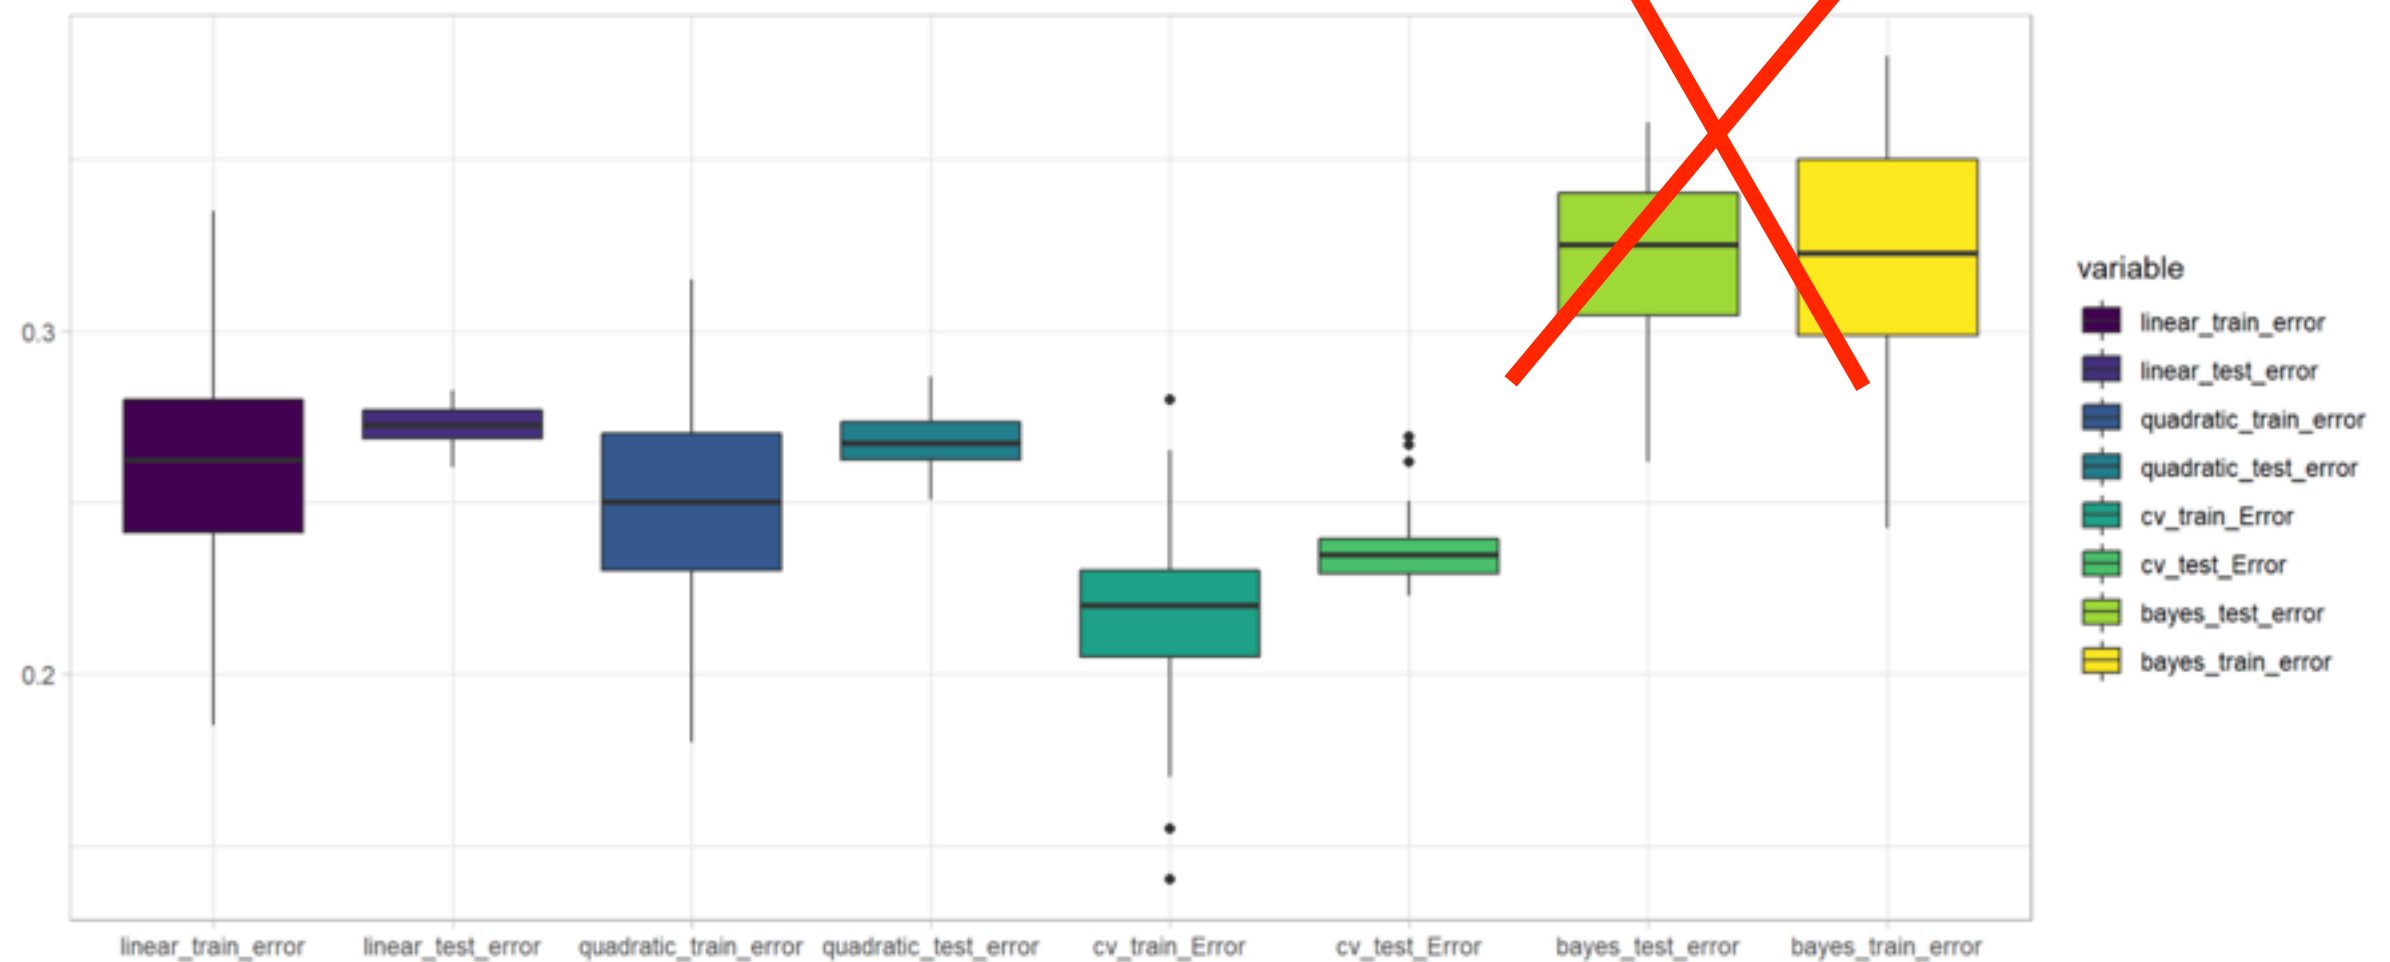
The three k-values, 1, 3, 5, used in Part I has nothing do with cvKNN in part II. No wonder your box plot does not look right.

```
def cvKNN(tr_x, tr_y, nfolds=10, k_vals=[1,3,5]):  
  
    pos_splits = list(split(list(np.where(tr_y == 1)[0]), nfolds))  
    neg_splits = list(split(list(np.where(tr_y == 0)[0]), nfolds))  
  
    cv_idxes = [pos_splits[i]+neg_splits[i] for i in range(nfolds)]
```





Apparently something is wrong.



```
Ytest_pred_Bayes = apply(testdata, 1, mixnorm, m1, m0, sqrt(1/5))  
Ytest_pred_Bayes = as.numeric(Ytest_pred_Bayes > 1/2)
```

**Apparently something is wrong.
The cut-off for mixnorm is NOT 1/2.**

Suggest to display the 8 sets of errors in ONE figure.

